L200921 Aisha Muhammad Nawaz BSCS 8A Midterm 2 Exam Solution MINING OF MASSIVE DATASETS SPRING 2024 DUE: 9th April 2024 (Tuesday)

```
#Running on Colab
!pip install pyspark
!pip install -U -q PyDrive
!apt install openjdk-8-jdk-headless -qq
import os
os.environ['JAVA_HOME'] = '/usr/lib/jvm/java-8-openjdk-amd64'
     Collecting pyspark
       Downloading pyspark-3.5.1.tar.gz (317.0 MB)
                                                     317.0/317.0 MB 3.0 MB/s eta 0:00:00
       Preparing metadata (setup.py) ... done
     Requirement already satisfied: py4j==0.10.9.7 in /usr/local/lib/python3.10/dist-packages (from pyspark) (0.10.9.7)
     Building wheels for collected packages: pyspark
       Building wheel for pyspark (setup.py) ... done
       Created wheel for pyspark: filename=pyspark-3.5.1-py2.py3-none-any.whl size=317488491 sha256=c4eb8e2cca258a0a74e9f717535dbbc1a38a276d72613b15bce96c8053231278
       Stored in directory: /root/.cache/pip/wheels/80/1d/60/2c256ed38dddce2fdd93be545214a63e02fbd8d74fb0b7f3a6
     Successfully built pyspark
     Installing collected packages: pyspark
     Successfully installed pyspark-3.5.1
     The following additional packages will be installed:
       libxtst6 openjdk-8-jre-headless
     Suggested packages:
       openjdk-8-demo openjdk-8-source libnss-mdns fonts-dejavu-extra fonts-nanum fonts-ipafont-gothic
       fonts-ipafont-mincho fonts-wqy-microhei fonts-wqy-zenhei fonts-indic
     The following NEW packages will be installed:
       libxtst6 openjdk-8-jdk-headless openjdk-8-jre-headless
     0 upgraded, 3 newly installed, 0 to remove and 45 not upgraded. Need to get 39.7 MB of archives.
     After this operation, 144 MB of additional disk space will be used.
     Selecting previously unselected package libxtst6:amd64.
     (Reading database ... 121753 files and directories currently installed.)
     Preparing to unpack .../libxtst6_2%3a1.2.3-1build4_amd64.deb ...
     Unpacking libxtst6:amd64 (2:1.2.3-1build4) ...
     Selecting previously unselected package openjdk-8-jre-headless:amd64.
     Preparing to unpack .../openjdk-8-jre-headless_8u402-ga-2ubuntu1~22.04_amd64.deb ...
     Unpacking openjdk-8-jre-headless:amd64 (8u402-ga-2ubuntu1~22.04) ...
     Selecting previously unselected package openjdk-8-jdk-headless:amd64.
     Preparing to unpack .../openjdk-8-jdk-headless_8u402-ga-2ubuntu1~22.04_amd64.deb ...
     Unpacking openjdk-8-jdk-headless:amd64 (8u402-ga-2ubuntu1~22.04) ...
     Setting up libxtst6:amd64 (2:1.2.3-1build4)
     Setting up openjdk-8-jre-headless:amd64 (8u402-ga-2ubuntu1~22.04) ...
update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/orbd to provide /usr/bin/orbd (orbd) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/servertool to provide /usr/bin/servertool (servertool) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/tnameserv to provide /usr/bin/tnameserv (tnameserv) in auto mode
     Setting up openjdk-8-jdk-headless:amd64 (8u402-ga-2ubuntu1~22.04) ..
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/clhsdb to provide /usr/bin/clhsdb (clhsdb) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/extcheck to provide /usr/bin/extcheck (extcheck) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/hsdb to provide /usr/bin/hsdb (hsdb) in auto mode update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/idlj to provide /usr/bin/idlj (idlj) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/javah to provide /usr/bin/javah (javah) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/jhat to provide /usr/bin/jhat (jhat) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/jsadebugd to provide /usr/bin/jsadebugd (jsadebugd) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/native2ascii to provide /usr/bin/native2ascii (native2ascii) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/schemagen to provide /usr/bin/schemagen (schemagen) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/wsgen to provide /usr/bin/wsgen (wsgen) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/wsimport to provide /usr/bin/wsimport) in auto mode
     update-alternatives: using /usr/lib/jvm/java-8-openjdk-amd64/bin/xjc to provide /usr/bin/xjc (xjc) in auto mode Processing triggers for libc-bin (2.35-0ubuntu3.4) ...
     /sbin/ldconfig.real: /usr/local/lib/libtbbmalloc.so.2 is not a symbolic link
     /sbin/ldconfig.real: /usr/local/lib/libtbbmalloc_proxy.so.2 is not a symbolic link
     /sbin/ldconfig.real: /usr/local/lib/libtbbbind_2_0.so.3 is not a symbolic link
     /sbin/ldconfig.real: /usr/local/lib/libtbbbind_2_5.so.3 is not a symbolic link
# Importing Required Libraries
import pyspark
from pyspark.sql import *
from pyspark.sql.functions import *
from pyspark import SparkContext, SparkConf
# Create Spark session and ContextRun PySpark.
# create the session
conf = SparkConf().set("spark.ui.port","4050")
# create the context
sc = pyspark.SparkContext(conf=conf)
spark = SparkSession.builder.appName("DataFrame").config('spark.ui.port', '4050').getOrCreate()
spark
```

```
SparkSession - in-memory
SparkContext
Spark UI
Version
v3.5.1
Master
local[*]
AppName
pyspark-shell
```

Q1 [marks]:

KMeans Algorithm is executed on a massive 2D dataset, and the result is saved in a file "F1". Your task is to input the huge result file F1 and write efficient SPARK code to compute the Davies-Bouldin Index (DBI) for evaluating the clustering quality.

The Davies-Bouldin index(DBI) measures the average similarity between each cluster and its most similar cluster, here, similarity is defined based on the Euclidean distance between cluster centroids.

```
Input File F1 format (2D point, cluster ID)
(1,2), C1
(5,5), C2
(2,2), C1
(11, 11), C3
```

Let's compute DB1 step by step using SPARK. You can use SPARK DataFrame or RDDs. Provide spark code for each part given below:

- a. Input File F1 in SPARK and compute the centroid of each cluster.
- b. Compute the average (avg) distance of points in each cluster to its centroid.
- c. Compute the similarity between each pair of clusters as follows:

For example, Similarity(Cluster 1, Cluster 2) = (Euclidean distance between centroid 1 and 2) / (maximum of (avg distances of points from centroids in C1, avg distances of points from centroids in C2)) Note: Avg distances are computed in part(b).

d. Compute the Davies-Bouldin Index(DBI) = Average of the pair-wise similarities between clusters computed in part(c).

```
[HINT] The Set of 2D points is huge(big data), but the number of clusters is very small compared to that.
import numpy as np
def getPointClusterPair(line):
     x,y,cluster=line.split(',')
x=float(x.replace('(',''))
      y=float(y.replace(')',''))
      cluster=cluster.replace(' ','')
      return cluster, (x,y,1)
# a. Input File F1 in SPARK and compute the centroid of each cluster.
f1=sc.textFile('F1.txt').map(lambda line: getPointClusterPair(line)) #Cluster as Key and value is point and count of point
#Summing up coordinates value in each cluster as well as the count of points
\label{lem:f1Centroid} f1Centroid=f1Sum.mapValues(lambda <math>x:(x[0]/x[2],x[1]/x[2]))
                                                                                                                                                                                                                           #Computing Avg by dividing sum of coordinates by count of coordinates in each cluster
f1Centroid.collect()
                [('C1', (1.5, 2.0)), ('C2', (5.0, 5.0)), ('C3', (11.0, 11.0))]
# b. Compute the average (avg) distance of points in each cluster to its centroid.
f1CentroidCollected=dict(f1Centroid.collect()) #Converting prev step ans to dictionary form
 f1DistCentroid = f1.map(lambda \ x: (x[0], ((x[1][0]-f1CentroidCollected[x[0]][0]) **2 + (x[1][1]-f1CentroidCollected[x[0]][1]) **2 + (x[1][1]-f1Centroid
f1D ist Centroid 2 = f1D ist Centroid. reduce By Key (lambda x, y: (x[0]+y[0], x[1]+y[1])) \\ \# Summing up dist and count of point in each cluster function of the point of t
f1DistCentroidFinal = f1DistCentroid2. mapValues(lambda \ x:(np.sqrt(x[\emptyset])/x[1])) \\ \qquad \# Computing \ Avg \ by \ dividing \ euc \ dist \ by \ count \ of \ coordinates \ in \ each \ clus \ for \ coordinates \ in \ each \ clus \ for \ coordinates \ in \ each \ clus \ for \ coordinates \ 
f1DistCentroidFinal.collect()
                [('C1', 0.3535533905932738), ('C2', 0.0), ('C3', 0.0)]
# c. Compute the similarity between each pair of clusters as follows:
# Similarity(Cluster 1, Cluster 2) = (Euclidean distance between centroid 1 and 2) / (maximum of (avg distances of points from centroids in C1,
     avg distances of points from centroids in C2)) Note: Avg distances are computed in part(b).
clusterDistAvg=dict(f1DistCentroidFinal.collect()) #Converting prev step ans to dictionary form
res={}
for cluster1,dist1 in clusterDistAvg.items():
      for cluster2,dist2 in clusterDistAvg.items():
            if(cluster1 != cluster2 and (cluster1+cluster2) not in res and (cluster2+cluster1) not in res): #To ensure no duplicate entries
                   maxDist=dist1
                   if(dist2>dist1):
                        maxDist=dist2 #Finding max dist
                   # Finding Euclidean Distance
                   euclideanDistance=np.sqrt((f1CentroidCollected[cluster1][0]-f1CentroidCollected[cluster2][0])**2+(f1CentroidCollected[cluster1][1]-f1CentroidCollected[cluster2][1])**2)
                   res[cluster1+cluster2]=euclideanDistance/maxDist if maxDist > 0 else 0 #Finally saving similarity in res dictionary for the pair
print(res)
                {'C1C2': 13.038404810405297, 'C1C3': 37.013511046643494, 'C2C3': 0}
```

4/8/24, 7:16 PM

d. Compute the Davies-Bouldin Index(DBI) = Average of the pair-wise similarities between clusters computed in part(c).
sumSim=0
count=0
for key,value in res.items():
sumSim=sumSim+value #Summing up all the similarites
 count=count+1 #Counting all the pairs

DBI=sumSim/count if count>0 else 0

Davies-Bouldin Index(DBI) = 16.683971952349598

print('Davies-Bouldin Index(DBI) = ',DBI)

Q2 [marks]: We applied user-based collaborative filtering on YouTube's video dataset. Each user gives a thumb-up or thumb-down sign on each video, which is changed to 1 and 0 in data.

V1 V2 V3 V4

UserA 1 0 1 1

UserB 0 1 1 1

- a) Find the similarity between User 1 and User 2 using Jaccard, Cosine, and Pearson Correlation.
- b) Which is the best measure for finding user similarity for the given dataset? Briefly explain your answer.
- c) In the case of Jaccard similarity, what should be done with missing values?

